

Keren Zhou

6100 Main ST – Houston, TX – 77005, United States
☎ +1-281-687-6961 ✉ keren.zhou@rice.edu 🌐 www.jokeren.tech

EDUCATION BACKGROUND

09/2017-07/2022	Department of Computer Science, Rice University Expected Degree: <i>Ph.D. in Computer Science</i> GPA: 3.9/4.0 Advisor: John Mellor-Crummey	Houston, United States
09/2014-07/2017	Institute of Computing Technology, Chinese Academy of Sciences Degree: <i>M.S. in Computer Architecture</i> GPA: 90/100 Advisor: Guangming Tan Thesis: High Performance Deep Learning Algorithms	Beijing, China
09/2010-07/2014	School of Software, Yunnan University Degree: <i>B.E. in Network Engineering</i> GPA: 92/100 (Rank: 1/290) Advisor: Wei Zhou Thesis: A Practical Concurrent Quadtree	Kunming, China

RESEARCH EXPERIENCE

09/2017-NOW	Rice University <i>Research Assistant</i> Scalable GPU Performance Measurement and Analysis Tool <ul style="list-style-type: none">Implemented OpenMP 5.0 OMPT Tool Interface for CUDA backend in llvm-openmp;Built a general runtime system to collect GPU activities on NVIDIA, AMD, and Intel GPUs and attributed them back to the corresponding CPU calling context;Analyzed GPU binaries to extract GPU functions, recover control flows, and map instructions to source code;Associated runtime samples with static GPU program structures to reconstruct calling context on GPUs and estimate instruction throughput and roof-line model;Studied HPC and machine learning applications, including PyTorch, Darknet, Quicksilver, Nekbone, Laghos, PeleC, QMCPACK, Nyx, GAMESS, NAMD, and LAMMPS. GPU Performance Advisor <ul style="list-style-type: none">Devised a method to attribute instruction stalls back to the instructions that caused them;Built a profile-guided performance advisor based on GPU performance metrics, program structures, and PC samples;Derived performance models to estimate speedups of individual suggestions proposed by the advisor;Optimized GPU applications by applying advice generated by the advisor to obtain speedups on a V100 GPU with 1.22× on average. GPU Value Redundancy Profiler <ul style="list-style-type: none">Investigated value redundancy problems in HPC and machine learning applications and achieved speedups by up to 1.93×.Built the first value profiler for NVIDIA GPUs to explore both temporal and spatial value redundancies in multi-node multi-GPU clusters;Devised asynchronous analysis and hierarchical sampling methods to reduce the tool overhead to 7.5× on average for Rodinia benchmarks.	Houston, United States
06/2015-07/2017	Institute of Computing Technology, Chinese Academy of Sciences <i>Research Assistant</i> GPU Performance Model <ul style="list-style-type: none">Decoded Nvidia GPU assembly codes and developed assemblers to generate GPU binaries;Built a static performance analysis model to estimate performance bottlenecks in GPU binaries. High Performance Deep Learning Framework <ul style="list-style-type: none">Devised a coarse-grained parallelism strategy with fine-grained vectorization and blocking effects on CPU, making CNNs 5-12 times faster than Caffe on a 16-core E5-2670;Wrote assembly codes to make use of dual issue and avoid bank conflict on GPU, improving convolution performance with up to 1.6× speedup than cuDNN on Kepler architectures.	Beijing, China

01/2013-07/2014 **Intelligent Web Laboratory, Yunnan University** **Kunming, China**
Research Assistant
Concurrent Data Structures

- Designed several concurrent multi-dimensional trees, including the first lock-free quadtree and k-d tree that are $1.09\times$ faster than state-of-the-art concurrent trees;

SELECTED PUBLICATIONS

JOURNALS.....

- [1] **Keren Zhou**, Guangming Tan, and Wei Zhou. Quadboost: A Scalable Concurrent Quadtree. In: *IEEE Transactions on Parallel and Distributed Systems* (TPDS), 2018

CONFERENCES.....

- [1] **Keren Zhou**, Xiaozhu Meng, Ryuichi Sai, John Mellor-Crummey. GPA: A GPU Performance Advisor Based on Instruction Sampling. In: *International Symposium on Code Generation and Optimization* (CGO), 2021
- [2] **Keren Zhou**, Yueming Hao, John Mellor-Crummey, Xiaozhu Meng, and Xu Liu. GVProf: A Value Profiler for GPU-based Clusters. In: *The International Conference for High Performance Computing, Networking, Storage, and Analysis* (SC), 2020
- [3] **Keren Zhou**, Mark Krentel, and John Mellor-Crummey. Tools for top-down performance analysis of GPU-accelerated applications. In: *34th ACM International Conference on Supercomputing* (ICS), 2020
- [4] **Keren Zhou**, Guangming Tan, Xiuxia Zhang, Chaowei Wang, and Ninghui Sun. A Performance Analysis Framework for Exploiting GPU Microarchitectural Capability. In *26th ACM International Conference on Supercomputing* (ICS), 2017
- [5] Xiuxia Zhang, Guangming Tan, Shuangbai Xue, Jiajia Li, **Keren Zhou**, and Mingyu Chen. Understanding GPU Microarchitecture to Achieve Bare-Metal Performance Tuning. In: *22nd ACM SIGPLAN Symposium on Principles and Practice of Parallel Programming* (PPOPP), 2017

POSTERS.....

- [1] **Keren Zhou**, Mark Krentel, and John Mellor-Crummey. A tool for top-down performance analysis of GPU-accelerated applications. In: *25th ACM SIGPLAN Symposium on Principles and Practice of Parallel Programming* (PPoPP), 2020
- [2] **Keren Zhou** and John Mellor-Crummey. A tool for performance analysis of GPU-accelerated applications. In: *International Symposium on Code Generation and Optimization* (CGO), 2019

INDUSTRY EXPERIENCE

- 05/2020-08/2020 **C++ Performance Optimization Team, Google Inc.** **Houston, United States**
Software Engineering Intern
 - Developed AutoDiff, a performance regression analysis tool that locates performance difference ranges, provide line and call site information, and allows flexible queries;
 - Proposed a method to recover high-resolution calling context with minimal overhead by augmenting call stacks for each instruction in LBR entries;
 - Reference: Software Engineer Wei Mi, wmi@google.com.
- 06/2018-08/2018 **PyTorch Team, Facebook Inc.** **Menlo Park, United States**
Research Intern
 - Accelerated neural networks on ARM CPUs using auto-tuning methods;
 - Analyzed Winograd algorithm's complexities of various convolution configurations;
 - Reference: Research Scientist Hao Lu, hlu@fb.com.
- 04/2017-07/2017 **Devtech Team, Nvidia Inc.** **Beijing, China**
Research Intern
 - Developed quantization tools on emerging GPUs to utilize INT8 capabilities;
 - Evaluated the precision and speed of different quantization modes on Pascal Titan X;
 - Reference: Technical Manager Julien Lai, julienlai@nvidia.com.

10/2013-02/2014 **Baidu Inc.** **Beijing, China**
Software Engineering Intern

- Optimized Hadoop workflow with its performance improved by 30%, making it capable of extracting thousands of features from raw text files and loading them into data warehouse;
- Developed a Hadoop workflow monitoring system that can display multiple workflow status and report exceptions;
- Reference: Senior Engineer Jing Li, lijing16@baidu.com.

AWARDS & HONORS

2020	ACM–IEEE-CS George Michael Memorial HPC Fellowship
2019	Ken Kennedy Institute ExxonMobil Fellowship
2019	Second Place, ACM CGO Student Research Competition
2017	Ken Kennedy Institute Andrew Ladd Fellowship
2017	Ken Kennedy Institute CS&E Fellowship
2016	Schlumberger Scholarship (3%)
2015	Top 10, Alibaba 1st Middleware Engineering Contest
2014	Bronze Medal, The 2014 ACM-ICPC Asia Regional Contest
2014	Outstanding B.E. Degree Thesis of Yunnan University
2013	Best Creative Award, Baidu Future Search Engine Contest
2013	Meritorious Winner, Mathematical Contest in Modeling
2011	Second Prize, China Undergraduate Mathematical Contest in Modeling
2011&2012&2016	National Scholarship (2%)

SELECTED PUBLIC TALKS

04/2021	Measurement and Analysis of GPU-accelerated Applications with HPCToolkit. NVIDIA GPU Technology Conference.
03/2021	Using HPCToolkit to Measure and Analyze the Performance of GPU-accelerated Applications. NERSC Invited Tutorial.
03/2021	GPA: A GPU Performance Advisor Based on Instruction Sampling. CGO'21.
11/2020	GVProf: A Value Profiler for GPU-Based Clusters. SC'20.
07/2020	Tools for Top-down Performance Analysis of GPU-Accelerated Applications. ICS'20.
03/2020	A Tool for Top-down Performance Analysis of GPU-accelerated Applications. PPOPP'20.
02/2020	Using HPCToolkit to Measure and Analyze the Performance of GPU-Accelerated Applications. ECP'20.
10/2019	Measurement and Analysis of GPU-computations Using HPCToolkit. BP Invited Talk.
08/2019	HPCToolkit—A tool for performance analysis for GPU-accelerated applications. Intel Performance Brown Bag.
08/2019	HPCToolkit + OpenMP. ECP/NERSC OpenMP Hackathon.
07/2019	Optimizing GPU-accelerated Applications with HPCToolkit. Scalable Tools Workshop.

ACADEMIC SERVICES

JPDC'21	Journal subreviewer
ICS'21	Conference paper subreviewer
ICDCS'21	Conference paper subreviewer
IPDPS'21	Conference paper subreviewer
PPoPP'21	Artifact Evaluation Committee member

SKILLS

Languages	C, C++, Java, Python, Bash, Go, JavaScript
Parallelism	Pthread, OpenMP, OpenCL, MPI, CUDA/HIP, DPCPP, RAJA/Kokkos